

Charles University in Prague
Faculty of Mathematics and Physics

MASTER THESIS



Bc. Jakub Kinšt

DEECo Cloudlets Exploratory Study

**Department of Distributed and Dependable Systems
Faculty of Mathematics and Physics, Charles University**

Supervisor of the master thesis: Doc. RNDr. Tomáš Bureš, PhD.

Study program: Informatics

Specialization: Software Systems

Prague 2015

TODO: thank supervisor and others

I declare that I carried out this master thesis independently, and only with the cited sources, literature and other professional sources.

I understand that my work relates to the rights and obligations under the Act No. 121/2000 Coll., the Copyright Act, as amended, in particular the fact that the Charles University in Prague has the right to conclude a license agreement on the use of this work as a school work pursuant to Section 60 paragraph 1 of the Copyright Act.

In Prague date

signature

Název práce: TODO: name in czech

Autor: Bc. Jakub Kinšt

Katedra / Ústav: Katedra distribuovaných a spolehlivých systémů

Vedoucí diplomové práce: Doc. RNDr. Tomáš Bureš, PhD.

Abstrakt: TODO: abstract

Klíčová slova: TODO: keywords

Title: DEECo Cloudlets Exploratory Study

Author: Bc. Jakub Kinšt

Department / Institute: Department of Distributed and Dependable Systems

Supervisor of the master thesis: Doc. RNDr. Tomáš Bureš, PhD., Department of Distributed and Dependable Systems, Faculty of Mathematics and Physics, Charles University

Abstract: TODO: abstract

Keywords: *TODO: keywords*

Contents

Contents	1
1 Introduction	3
1.1 Motivation	3
1.2 Thesis goals	4
1.3 Problems and challenges	4
2 Background	5
2.1 DEECo component model	5
2.1.1 Ensemble-based component systems	5
2.1.2 Running example	6
2.1.3 Basic principles	6
2.1.3.1 Component	6
2.1.3.2 Ensemble	8
2.1.4 jDEECo library	8
2.1.5 CDEECo library	10
2.2 Concept of Cloudlets	10
2.3 Computation offloading	11
2.3.1 Use cases	11
2.4 Android platform fundamentals	12
2.4.1 Application development	12
2.4.1.1 AndroidManifest.xml file	12
2.4.1.2 Permissions	14
2.4.1.3 Activities	14
2.4.1.4 Services	14
2.4.1.5 Development, building and running applications	16
3 Analysis	17
3.1 Offloading controlled by EBCS	17
3.1.1 Proposed adaptation architecture	18
3.2 Towards the reference architecture	20

3.2.1	Suitable DEEC _o implementation	20
3.2.2	Offloadable application components	20
3.2.3	Stateful components	20
3.2.4	Deployment strategies	21
3.2.5	Communication between layers	22
3.2.6	Framework API	22
3.2.7	Java and Android together	22
4	Reference Architecture	23
4.0.8	jDEEC _o on Android and UDP (TODO)	23
4.0.9	Offloadable app components	23
4.0.10	DEEC _o control layer	23
4.0.11	Putting it all together	23
5	Implementing the Reference Architecture	24
6	Demonstration Applications	25
7	Making an Android Application Offloadable	26
8	Conclusion	27
9	Attachments	28
	List of Abbreviations	29
	List of Figures	30
	Bibliography	31

Chapter 1

Introduction

1.1 Motivation

With growing usage of mobile devices such as smartphones, tablets and wearables, users tend to use these devices to do things they used to do on desktop computers or laptops. Unlike desktop computers or laptops, these devices are usually not connected to a power source when their users are using them. The computation power of these devices is growing much more rapidly than the capacities of used batteries. With hardware being more and more efficient, the battery life of such devices is relatively stable for past few years, but still, it is very limited. This means that when developer is designing a mobile application, it is essential to always think of mobile-specific resources that the application uses and design the application to be as efficient as possible.

When developing applications performing resource-intensive computation tasks such as image recognizing, sound recognizing or even communicating with a remote API, there is one technique that can be particularly useful to save at least some of the resources - offloading. Offloading resource-intensive tasks to another machine can sometimes significantly reduce resource usage when used wisely.

Such devices also usually form a very dynamic large-scale environment. As their users are leaving their homes, entering workplaces, using WiFi networks during transportation, in coffee shops, restaurants, their smartphones or tablets are entering and leaving different networks very often. This dynamic environment tend to be very hard to control. With nodes appearing and disappearing on regular basis it is necessary to design a system without a dedicated central authority to control such environments.

Recently introduced Ensemble-Based Component Systems (EBCS) could provide an optimal solution for this kind of problem since they are specifically tailored for dynamic large-scale decentralized environment. EBCS could be used

as a control layer in our offloading scenario.

1.2 Thesis goals

The goal of this master thesis is to explore possibilities of using the DEEC_o component model to achieve cloudlet offloading of parts of a mobile application. This means executing parts of application computation on a different machine within local network.

First of all it is necessary to port the existing Java implementation of DEEC_o component model to work on the Android platform to be able to run it's runtime on mobile devices. The existing library does not support communication between nodes in "real-life" computer networks. Therefore the next step is to develop an extension to the library providing support for a communication within local network.

The most important goal is to design and implement a reference architecture which brings offloading capabilities to a regular mobile application for the Android platform. This architecture should use the principles proposed in the DEEC_o component model to create a control layer for the offloading mechanism. This goal also includes choosing as universal way as possible to break the application into separate components that are prepared to be offloaded - run on a different machine. Also a mechanism of deciding which offloadable component will run on which machine in the network.

Last but not least, a demonstration application is to be developed to bring the idea to the "real world" and show the offloading mechanism in action. The demo application should provide a reasonable functionality that could benefit from being offloaded off the mobile device itself and potentially save some of it's battery life or other limited resource. The developed application should be evaluated in terms of saved resources while offloading is enabled. The application (together with a set of instruction) should serve as an example of the process of adding offloadable capabilities to any other Android application as easily as possible.

1.3 Problems and challenges

TODO

Chapter 2

Background

2.1 DEECo component model

2.1.1 Ensemble-based component systems

With recent increasing possibilities provided by the evolution in the field of mobile devices and their improving connectivity, new ways of addressing social and environmental challenges (ambient assisted living, smart city infrastructures, emergency coordination, environmental monitoring) are emerging. Solutions are achieved by building large-scale Resilient Distributed Systems (RDS) that respond to and influence activities in the real world. “As RDS have to cope with very dynamic and open-ended environments, they exhibit a high degree of adaptivity and autonomicity.”[7]

The dynamic and autonomic nature of RDS causes issues with scalability when using traditional component architectures and models. Ensemble-based component systems (EBCS) were introduced to overcome this problem in [7]. The key feature of EBCS is the different composition of components. Instead of explicit component architecture, components are implicitly formed into so-called *ensembles* based on declared predicates which makes the composition of components very dynamic and adaptive.

Finally, the DEECo (Distributed Emergent Ensembles of Components) component model was proposed “to refine the principles of EBCS into a systematic approach for building software for RDS”[7]. DEECo is an instance of EBCS which comes with a framework for building applications benefiting from EBCS features described earlier.

2.1.2 Running example

An example provided in [7] is very useful for highlighting the key features of EBCSs. It is based on the electrical vehicle navigation case study.

The case study involves e-vehicles and their navigation around a city. Each driver of one of the e-vehicles has his own plan of stops (POIs) in a form of calendar. The vehicles can only park in special parking places equipped with a charger clustered in parking lot/charging stations (PLCS). They are also able to constantly monitor their position, energy consumption and battery level. The vehicles can communicate with each other as well as with the parking/charging stations to be able to plan/reserve a parking place. The objective of the case study is to coordinate journey planning with constraints coming from the parking/charging strategy.

The important factor is that no central coordination authority is assumed in this case study. It involves highly dynamic environment of multiple nodes. The architecture of those nodes changes constantly in time based on the current position of e-vehicles. These are exactly the challenges targeted by EBCS.

2.1.3 Basic principles

DEECo is based on two main concepts - *component* and *ensemble*. According to [7], a component is an independent, self-sustained unit of development, deployment and computation. An ensemble, on the other hand, is a dynamic binding mechanism linking a set of components together and providing a communication channel between them. Actually, one of the most important principle of DEECo is that ensembles provide the only way of communication between components. On top of components and ensembles, DEECo defines a runtime, which provides necessary management services for both.

2.1.3.1 Component

A single component consists of *knowledge* and *processes*. Knowledge reflects the state of the component and it is organized in a hierarchical data structure mapping knowledge identifiers to values. It is exposed through an implicit set of interfaces which means the interface represents a partial view of the knowledge.

As far as the processes are concerned, they are basically tasks, which are able to manipulate the component's knowledge. They are represented by a function which has a set of input and output knowledge fields. The function is called by the DEECo runtime framework and the process of reading the input knowledge, running the function and writing the output knowledge is atomic. A component should never communicate with other components directly via processes. It

```

interface AvailabilityAggregator:
    calendar, availabilities

interface AvailabilityAwareParkingLot:
    position, availability

component Vehicle features AvailabilityAggregator:
    knowledge:
        batteryLevel = 90%,
        position = GPS(...),
        calendar = [ POI(WORKPLACE, 9AM-1PM), POI(MALL, 2PM-3PM)
                    , ... ],
        availabilities = [ ],
        plan = {
            route = ROUTE(...),
            isFeasible = TRUE
        }
    process computePlan:
        in plan.isFeasible, in availabilities, in calendar,
        inout plan.route
        function:
            if (!plan.isFeasible)
                plan.route <- Planner.computePlan(calendar,
                    availabilities)
            scheduling: triggered( changed(plan.isFeasible) OR
                changed(availabilities) )
    process checkPlanFeasibility:
        in plan.route, in batteryLevel, in position, out plan.
        isFeasible
        function:
            plan.isFeasible <- Planner.isFeasible(plan.route,
                batteryLevel, position)
            scheduling: periodic( 5000ms )

component PLCS features AvailabilityAwareParkingLot:
    knowledge:
        position = GPS(...),
        availability = ...
    process observeAvailability:
        out availability
        function:
            availability <- Sensors.getCurrentAvailability()
            scheduling: periodic( 2000ms )

```

Figure 2.1: Example of DEEC_o component definitions in a DSL (Source: [7])

```

ensemble UpdateAvailabilityInformation:
  coordinator: AvailabilityAggregator
  member: AvailabilityAwareParkingLot
  membership:
    if poi in coordinator.calendar:
      distance(member.position, poi.position) <= TRESHOLD
      && isAvailable(poi, member.availability)
  knowledge exchange:
    coordinator.availabilities <- { (m.id, m.availability) |
      m in members }
  scheduling: periodic( 5000ms )

```

Figure 2.2: Example of DEECo ensemble definition in a DSL (Source: [7])

should only read and/or manipulate knowledge of it's own. The process can be either triggered by a certain event, or it can be performed periodically by the runtime. See Figure 2.1 for an example of component definitions.

2.1.3.2 Ensemble

An ensemble serves as a binding mechanism between multiple components. It provides a way of communication between them. One of the components involved in an ensemble is in the position of *coordinator* and others are just *members*. The involvement of components in an ensemble is defined by ensemble's *membership* function. The membership function consists of the definition of the *coordinator interface*, the *member interface* and the *membership predicate* which is evaluated by the DEECo runtime to determine which two components represent a coordinator-member pair.

The communication between the components within an ensemble is provided through the *knowledge exchange* function. Actually, the function provides interaction between the coordinator and all other members only, so all communication must go through the coordinator. Communication is realized by the ability to reading and/or writing a knowledge of both the coordinator and the member within the coordinator and member interfaces. The knowledge exchange operation can be either triggered by a certain event, or it can be performed periodically by the runtime. See Figure 2.2 for an example of an ensemble definition.

2.1.4 jDEECo library

To be able to actually use the DEECo component model for development of RDS, a framework called jDEECo[9] has been developed. It is an implementation of the model, which provides a way to deploy and run real DEECo-based applications

```

@Component
public class Vehicle {
    public List<CalendarEvent> calendar;
    public Plan plan;
    public EnergyLevel batteryLevel;
    public Map<ID, Availability> availabilities;
    public Position position;

    public Vehicle() {
        // initialize the initial knowledge structure reflected
        // by the class fields
    }

    @Process
    public static void computePlan(
        @In("plan.isFeasible") @Triggered Boolean isPlanFeasible
        ,
        @In("availabilities") @Triggered Map<...>availabilities,
        @In("calendar") List<CalendarEvent> calendar,
        @InOut("plan.route") Route plannedRoute
    ) {
        // re-compute the vehicle's plan if it's infeasible
    }

    @Process
    @PeriodicScheduling(5000)
    public static void checkPlanFeasibility(
        @In("plan.route") Route plannedRoute,
        @In("batteryLevel") EnergyLevel batteryLevel,
        @In("position") Position position,
        @Out("plan.isFeasible") OutWrapper<Boolean>
            isPlanFeasible
    ) {
        // determine feasibility of the plan
    }
    ...
}

```

Figure 2.3: Example of jDEECo component (Source: [7])

written in Java language. A jDEECo component has a form of a Java class marked by the `@Component` annotation. Its knowledge is represented by public non-static fields and the processes are defined by public static methods marked by the `@Process` annotation. Its input and output knowledge fields are marked by the `@In`, `@Out` or `@InOut` annotations.

Similarly, a jDEECo ensemble is defined by annotating a standard Java class with the `@Ensemble` annotation. The membership and knowledge exchange functions are defined by public static methods marked by `@Membership` or `@KnowledgeExchange` respectively. Its input and output knowledge fields should be marked in the same manner as with the component's process functions. An example of jDEECo component definition is provided in Figure 2.3.

2.1.5 CDEECo library

Recently, a C++ realization of the DEECo model has been developed to support DEECo application development on different platforms where Java environment is not available. It is called CDEECo[8].

2.2 Concept of Cloudlets

A *cloudlet* is an emerging architectural element arising from the combination of mobile computing and cloud computing. [4, 15] It is a middle layer of recently proposed 3-layer hierarchy, which consists of a mobile device, cloudlet and cloud as we know it. The goal of the cloudlet proposal is to “bring the cloud closer” to the mobile device. According to [4], a cloudlet features four key attributes:

- **soft state only** - A cloudlet should not have any hard state, but it can contain a cached state from the cloud for instance. The lack of a hard state makes cloudlets self-managing and easy to deploy.
- **powerful, well-connected and safe** - A cloudlet should feature sufficient compute power compared to mobile device. It should have an unlimited power source (connected to a power outlet) and very good and stable connectivity to the cloud.
- **close at hand** - It should be close to the mobile device in terms of latency and high bandwidth. Usually this means connected on the same local network via Wi-Fi.
- **builds on standard cloud technology** - It should be similar to classic cloud infrastructures.

The concept of cloudlets was designed for mobile application computation offloading scenario.[15] According to this paper, cloudlets are the technology bringing a new type of mobile applications which are resource-intense but latency-sensitive at the same time. These applications are expected to emerge in the near future.

2.3 Computation offloading

Computation offloading generally means delegating certain computing tasks to another node - usually a cloud, cluster or grid. The goal of computation offloading is either to save system resources used on a device or to access higher computational performance unavailable when performing the task on its own.

In the world of smartphones, tablets or wearables, developers are facing new challenges and problems such as limited battery life and lower performance. When developing a mobile application for such devices, transferring resource-intensive tasks to the cloud can significantly reduce battery usage and bring better user experience to the user.[12]

2.3.1 Use cases

Computation offloading is starting to be widely used strategy in certain areas of mobile applications. For instance, each major mobile platform now has its own “voice assistant” (Android: *Google Now*, iOS: *Siri*, Windows Phone: *Cortana*) which is able to listen for whatever question the user may possibly have and respond instantly with an answer. [14] Let’s focus on a voice search for a contact name present on the mobile device for instance. This process consists of two main steps: (i) speech recognition which converts a recording of human voice to a search query and (ii) actual search for the contact based on the name. If the process of voice recognition was performed locally on the mobile device itself, the response would most definitely be longer than instant. That is because voice recognition is highly non-trivial task and mobile devices usually don’t feature such computational power to provide response in a reasonable time. Here is what actually happens (after simplification): (i) the mobile application preprocesses the voice recording provided by the user, (ii) sends the preprocessed recording to the cloud and (iii) receives a string representation of the search query which was computed at a datacenter with enormous computation power.[14] The search query can then be used for a simple local search task.

However, computation offloading is definitely not a silver bullet. Since network communication used to reach the offloading target is considered resource-intensive as well, one has to find an optimal balance when designing application

featuring computation offloading. It is necessary to carefully inspect if the nature of a task is suitable for offloading. Generally, most resources can be spared by offloading long-running tasks or by batching (if possible) multiple operations and offloading it at once. A very good examples of such long-running tasks are those involving some kind of recognition - voice recognition, image recognition, song-recognition, augmented reality, video recognition and many more.[10, 13, 11, 12, 14]

2.4 Android platform fundamentals

Android[3] is an operating system by Google for mobile devices such as smart-phones, tablet computers, smartwatches and other. It is based on the Linux kernel and the apps are written primarily in Java using the Android SDK[1].

Android does not use Java Virtual Machine as one would expect, but instead, it uses *Dalvik* with just-in-time (JIT) compilation as a virtual machine to run the applications. Since Android 5.0 Lollipop, *Dalvik* was replaced by *ART*, which introduced ahead-of-time (AOT) compilation. The compilation is performed at the time of the installation of an application.

Also the APIs of both platforms differ. Most of the standard API is the same, but many differences are present. For instance the classes related to *AWT* or *Swing* are obviously not present, since Android uses it's own user interface implementation. Also vendor packages like `com.sun.*` are unavailable. On the other hand, the Android API includes other classes specific to the platform.

These differences are usually sources of issues when developing a library (or any code at all), which should work on both Java SE and Android. Therefore the developer should have these differences on mind at all times during development.

2.4.1 Application development

2.4.1.1 AndroidManifest.xml file

Every Android application must contain a special XML file which describes the application for the operating system. It contains information about application package - universal application identifier, it's name, icon, permission requests, hardware requirements and more. It also contains definitions of main application building blocks such as *Activities*, *Services*, *BroadcastReceivers* and others.

An example of *AndroidManifest.xml* file could look like this:

```
<?xml version="1.0" encoding="utf-8"?>
<manifest xmlns:android="http://schemas.android.com/apk/res/
    android"
    xmlns:tools="http://schemas.android.com/tools"
```



```

package="cz.kinst.jakub.example.app">

<uses-permission android:name="android.permission.INTERNET"/>
<uses-permission android:name="android.permission.GET_ACCOUNTS"
"/>
<uses-permission android:name="android.permission.
    READ_EXTERNAL_STORAGE"/>
<uses-permission android:name="android.permission.
    WRITE_EXTERNAL_STORAGE"/>

<uses-feature android:name="android.hardware.camera"
    android:required="true"/>

<application android:name=".ExampleApplication"
    android:icon="@drawable/ic_launcher"
    android:label="@string/app_name"
    android:theme="@style/AsparagusTheme">

    <activity android:name=".ui.activities.MainActivity"
        android:icon="@drawable/ic_launcher"
        android:logo="@drawable/logo_font"
        android:label="@string/app_name">
        <!-- This activity is the default activity listed in the
            app drawer on the device. It provides an entry point
            for the user. -->
        <intent-filter>
            <action android:name="android.intent.action.MAIN"/>
            <category android:name="android.intent.category.LAUNCHER"
                "/>
        </intent-filter>
    </activity>
    <activity android:name=".ui.activities.SettingsActivity"
        android:icon="@drawable/ic_asparagus"
        android:label="@string/settings">
    </activity>
    <receiver android:name=".sync.GcmBroadcastReceiver"
        android:permission="com.google.android.c2dm.permission.
            SEND">
        <intent-filter>
            <action android:name="com.google.android.c2dm.intent.
                RECEIVE"/>
            <category android:name="cz.kinst.jakub.example.app"/>
        </intent-filter>
    </receiver>
    <service android:name=".sync.SyncAdapterService"
        android:exported="true">
        <intent-filter>
            <action android:name="android.content.SyncAdapter"/>
        </intent-filter>

```

```
        <meta-data android:name="android.content.SyncAdapter"
            android:resource="@xml/sync_adapter"/>
    </service>
</application>
</manifest>
```

2.4.1.2 Permissions

On Android platform, any operation that could be potentially considered as dangerous for the user must be declared in the *AndroidManifest.xml* file statically. These operations include access to the Internet, reading writing to the external storage, access to contact list, sending a text message, placing a call and many others.

When installing an application all requested permissions are presented to the user for acceptance. Since then, the application is granted all of the permissions and is able to perform desired operations falling under those permissions. See the example of the *AndroidManifest.xml* file above for examples of requesting permissions.

2.4.1.3 Activities

An activity is probably the most important building block of every Android application. It is a visual component represented by a “window” (usually full-screen if not specified otherwise) and it is used to display user interface (UI). One or more activities can be tagged as launchable from application list (see the example of the *AndroidManifest.xml* file above) and others can be started manually using *Intents*.

An activity is represented by a Java class, which describes its functionality. The structure (layout) of the user interface is described in separate reusable XML files using proprietary markup.

Android applications are started differently than standard Java applications. There is no `main()` method which is run automatically. Instead, when the activity is started, special callback methods are called according to standard activity lifecycle.

2.4.1.4 Services

Services are non-visual components providing a way to run arbitrary code in the background without any user interface. Their lifecycle is similar to activity lifecycle displayed on figure 2.4. A service can be started via *Intents* but it can be also set up to be invoked regularly by an *AlarmManager* etc.

2.4.1.5 Development, building and running applications

The recommended development environment for developing Android applications is currently the *Android Studio*, which is based on *IntelliJ IDEA* from JetBrains. The apps should be built using *Gradle*^[2] build system, which is relatively new and it's Android support is still in development.

Android Studio together with *Gradle* takes care of building and packaging the application. The output is an *.apk* file containing entire application. This file can be transferred to the device and installed simply by launching it. The developer can also run the application directly on connected device from the IDE.

Chapter 3

Analysis

3.1 Offloading controlled by EBCS

Using EBCS as a control layer for computation offloading of mobile applications was proposed in [5]. Particularly, this paper focuses on the management of ad-hoc cloud (cloudlet) systems in such dynamic environment without any need for central authority and with focus on scalability and robustness of the solution.

The paper provides us with a running example which contains a user with a tablet computer traveling in a train or a bus, who wants to do productive work during the transportation. When the user entered the train/bus, the tablet automatically registers there is an offload server available connected to the same Wi-Fi network. In order to save battery, the tablet offloads most resource-intensive tasks to that server. When the train/bus reaches its destination, the server informs the tablet that it will soon be unavailable so tasks will start to move back to local execution. However, the tablet then registers another offload server. This time, it is provided by the train/bus terminal authority. Again, the tablet can move some of its tasks to this node.

The paper then generalizes the idea assuming a mobile device **M** and two stationary devices **S** and **T** (offload servers). **M** runs application **A** which can be divided to **A_f** and **A_b** - a frontend responsible for user interaction and backend responsible for computing resource-intensive tasks.

The scenario from the running example can be generalized to following:

1. **M** discovers **S**
2. **M** assesses that it is a better alternative to offload **A_b** to **S** in regards to battery usage
3. **S** notifies **M** that it will be soon unavailable




Figure 3.1: Adaptation architecture of the running example: phases 1 (M isolated), 2 (S discovered), and 3 (Ab migrated to S). Phases 1,2,3 are in the figure denoted by (1), (2), (3). (Source: [5])

4. **M** discovers **T**

5. **M** assesses that it is a better alternative to offload **Ab** to **T** in regards to battery usage

The challenge is “in predicting which deployment scenario will—in the context of ad-hoc cloud—deliver the expected user experience”. [5] The paper also assumes that each offloadable part will have certain performance model, which provides a rough estimate of user experience for each deployment plan. Parts of the application can then be deployed dynamically based on the predicted user experience when running on different nodes.

3.1.1 Proposed adaptation architecture

Architecture proposed in the paper is designed to form a separate layer mirroring the architecture of the adapted application in order to separate concerns as much as possible. The architecture consists of a set of components and a set of ensembles linking the components together.

Planner component Each adapted application is represented by one Planner component. This component is responsible for deciding a deployment plan based on provided alternatives for deployment. It basically means comparing all possible scenarios of backend deployment and choosing which backend component will run on which device. The input of this operation is NFP-related data (*NFPData*) containing measured estimate of potential performance of particular backend deployment.

Monitor component Each offloadable part of application (backend application component in our case) is mirrored by a set of Monitor components representing its deployment alternatives on different devices. It is responsible for calculating *NFPData* for the particular alternative. The Monitor is either in *running mode*, which means that the corresponding application component is currently running on the node this Monitor belongs to or in *mock mode* - the application component is running on different node. Based on the current mode, the *NFPData* is obtained by measuring of running component or estimated by provided performance measurement model based on the machine-specific performance data (*NFPDeviceData*)

Device component Device component mirrors each computation node and it is responsible for spawning new Monitor components for each application component based on *MonitorDefs* provided by the Planner.

Planner-Device ensemble To inform Device components about application components available for offloading, it must be involved in an ensemble together with the Planner. Planner is the coordinator of the ensemble and monitor definitions (application component definitions) are being pushed to Device components. This is the point where the deployment can be constrained - it can be defined which application components can be deployed on which nodes.

Planner-Monitor ensemble This ensemble's purpose is to gather *NFPData* from all available Monitors and push them to the Planner so it has all possible alternatives available for planning the deployment strategy.

Device-Monitor ensemble Each device pushes its *NFPDeviceData* to its Monitors in mock mode through this ensemble to let them estimate potential performance of the alternative they correspond to.

The number of nodes and Monitors can change very dynamically, since the devices are appearing and disappearing constantly. This is a great opportunity to use the emergent component ensembles, because they were proposed exactly for this type of situations. Thanks to exploiting the features of EBCS, the adaptation architecture is scalable and robust. The architecture is also very flexible - the deployment can be arbitrarily constrained, the *NFPData* can be obtained by a measurement or by a performance model. According to [5], it is also scalable for extensions: "For instance the Planner itself can be subject to migration in case the application does not have any frontend. Additionally, when understanding the Planner as an entity controlling the NFPs of the application, it is possible to foresee the existence of multiple Planners per application, thus hierarchically decomposing the adaptation."

3.2 Towards the reference architecture

Since one of the goals of this thesis is to design and implement a reference architecture for offloading mobile applications with DEEC_o as the control layer and DEEC_o is the implementation of EBCS, we can build the reference architecture based on the adaptation architecture described above in Section 3.1. On the way to the solution, few design choices have to be made and some challenges have to be dealt with. In this Section we will discuss them.

3.2.1 Suitable DEEC_o implementation

The first challenge lies in proper DEEC_o implementation. As stated before, two implementations have been developed up until now. A Java implementation called jDEEC_o and C++ implementation named CDEEC_o. The goal is to implement a solution for Android mobile platform, the obvious choice is to use the Java implementation.

The first issue is that, as discussed above, not all standard Java classes are available on the Android platform and it will most likely be necessary to perform certain changes to the framework to make it compatible with Android.

Other big challenge is introduced by the fact that jDEEC_o has practically never been used in a real-life application. It was tested in many simulations, but the framework lacks any support for network communication between nodes. Therefore, an extension providing such capability has to be developed for the jDEEC_o library.

3.2.2 Offloadable application components

To enable parts of an application to be deployed on different machines, one must design an universal, compact and scalable solution for splitting an application into offloadable components. Such components should be encapsulated and provide a common interface for interaction over a wireless network.

Regarding how the components should be deployed on the nodes, there are probably two main options to choose from. Either the code is installed on all possible nodes in advance statically, or some kind of source code/binary migration has to be used. Latter would enable dynamic component deployment to new nodes.

3.2.3 Stateful components

Many use cases for computation offloading involve some kind of a soft state (perhaps even a hard state). A voice recognition component may feature a

history of queries improving further recognition, a text recognition component may have some kind of basic user preferences, components may use a caching mechanism to improve their performance and many more.

Such feature definitely adds certain amount of complexity to the solution, because in this kind of dynamic environment, nodes appear and disappear on regular basis. Often they disappear suddenly and without any previous notice. (Wi-Fi out of range, Wi-Fi disabled by the user, etc.) The solution should therefore include a mechanism addressing this type of situation by seamlessly transferring the state between the nodes. The mechanism should be as robust as possible (given circumstances).

3.2.4 Deployment strategies

As described in the adaptation architecture, when deciding current deployment plan for all the components, many strategies are available for consideration. Deployment strategy generally consists of two steps: (i) obtaining *NFPData* (an estimate of the user experience) for deployment of particular component on certain node and (ii) deciding, based on provided set of *NFPData* for each component and each node, which application component will run on which node in the new deployment plan. Both steps can be customized in many ways.

In case the component is currently deployed on a different node, one way of calculating *NFPData* is using provided performance model involving device-specific data - for example a function of currently available memory, CPU usage, battery level, connectivity quality, or any other metric available - to obtain a rough estimate of user experience provided by corresponding alternative. In some cases it might be better to actually measure the execution of the component with a small static sample input. For instance, when implementing a text-recognition component, the performance can be calculated by measuring the execution time of recognizing a short word like “apple”. However, not all use cases provide this option. On the other hand, when the component is currently deployed on the corresponding node, one can simply measure the execution based on some metric. In our text-recognizing scenario, we could for example measure how long in average does it take to recognize one word. The other option is not to take the current mode into account and simply use the performance metric or small sample execution in both cases.

Later, when deciding which deployment alternative is the best for certain application component, a function choosing the best *NFPData* value from provided set has to be designed. This is the second entry point for potential customizations.

The reference architecture in the form of implemented framework should provide an easy way for customizing deployment strategies suitable for each use

case.

3.2.5 Communication between layers

Mentioned in the jDEECo library description, the process methods of DEECo components are executed in a static context by the jDEECo runtime. Therefore no reference to the component itself or its dependencies is available inside these methods. However, the components need to have some channel for communication with a “world outside DEECo” - the components at the application layer need to be notified when DEECo control layer decides to change the deployment strategy for instance. This channel has to be designed and implemented as well.

3.2.6 Framework API

One of the goals is making the developed framework easy-to-use for transforming a regular Android application to an offloadable application where offloading is managed by DEECo. This involves designing a straightforward interface for smooth transition which might be quite challenging to do. The API should interfere with existing application’s codebase as little as possible.

3.2.7 Java and Android together

The reference architecture implementation should serve as a framework to make standard Android apps offloadable. The developed framework has to be compatible with Android platform as well as the desktop (Java SE/EE) since the purpose of the work is to enable parts of Android applications to be executed on more powerful nodes (Servers, Desktops). This requirement can bring certain challenges with compatibility that have to be considered during the development. The code must be either compatible with both platforms or split into two branches for either platform with as much common code as possible.

Chapter 4

Reference Architecture

4.0.8 jDEECoo on Android and UDP (TODO)

TODO:

4.0.9 Offloadable app components

TODO:

4.0.10 DEECoo control layer

TODO:

4.0.11 Putting it all together

TODO:

Chapter 5

Implementing the Reference Architecture

Chapter 6

Demonstration Applications

Chapter 7

Making an Android Application Offloadable

Chapter 8

Conclusion

Chapter 9

Attachments

TODO: attached DVD

List of Abbreviations

AOT	Ahead-of-time
API	Application Programming Interface
APK	Android application package
ART	Android Runtime
DEECo	Dependable Emergent Ensembles of Components
EBCS	Ensemble-based component systems
IDE	Integrated Development Environment
JIT	Just-in-time
PLCS	Parking Lot/Charging Station
POIs	Points Of Interest
RDS	Resilient Distributed Systems
SDK	Software Development Kit
UI	User Interface
XML	Extensible Markup Language

List of Figures

2.1	Example of DEEC _o component definitions in a DSL (Source: [7])	7
2.2	Example of DEEC _o ensemble definition in a DSL (Source: [7]) . .	8
2.3	Example of jDEEC _o component (Source: [7])	9
2.4	Android activity lifecycle diagram. (Source: http://developer.android.com/reference/android/app/Activity.html)	15
3.1	Adaptation architecture of the running example: phases 1 (M isolated), 2 (S discovered), and 3 (Ab migrated to S). Phases 1,2,3 are in the figure denoted by (1), (2), (3). (Source: [5])	18

Bibliography

- [1] Android SDK. <http://developer.android.com>. Accessed: 2015-02-25.
- [2] Android Tools: Gradle Plugin User Guide. <http://tools.android.com/tech-docs/new-build-system/user-guide>. Accessed: 2015-02-25.
- [3] Android™. <http://android.com>. Accessed: 2015-02-25.
- [4] Elijah: Cloudlet-based Mobile Computing. <http://elijah.cs.cmu.edu/>. Accessed: 2015-03-13.
- [5] Lubomir Bulej, Tomas Bures, Vojtech Horky, and Jaroslav Keznikl. Adaptive deployment in ad-hoc systems using emergent component ensembles: Vision paper. In *Proceedings of the 4th ACM/SPEC International Conference on Performance Engineering*, ICPE '13, pages 343–346, New York, NY, USA, 2013. ACM.
- [6] Gerostathopoulos I. Hnetyinka P. Keznikl J. Kit M. Plasil F. Bures, T. Autonomous components in dynamic environments. *Awareness: Self-Awareness in Autonomic Systems Magazine*, pages 249–252, September 2012. <http://www.awareness-mag.eu/pdf/004415/004415.pdf>.
- [7] Tomas Bures, Ilias Gerostathopoulos, Petr Hnetyinka, Jaroslav Keznikl, Michal Kit, and Frantisek Plasil. Deeco: An ensemble-based component system. In *Proceedings of the 16th International ACM Sigsoft Symposium on Component-based Software Engineering*, CBSE '13, pages 81–90, New York, NY, USA, 2013. ACM.
- [8] D3S, Charles University in Prague. CDEECo website. <https://github.com/d3scomp/JDEECo>. Accessed: 2015-02-25.
- [9] D3S, Charles University in Prague. jDEECo website. <https://github.com/d3scomp/JDEECo>. Accessed: 2015-02-25.

- [10] Dejan Kovachev and Ralf Klamma. Framework for computation offloading in mobile cloud computing. *International Journal of Interactive Multimedia and Artificial Intelligence*, 1(7):6–15, 2012.
- [11] K. Kumar and Yung-Hsiang Lu. Cloud computing for mobile users: Can offloading computation save energy? *Computer*, 43(4):51–56, April 2010.
- [12] Zhiyuan Li, Cheng Wang, and Rong Xu. Computation offloading to save energy on handheld devices: A partition scheme. In *Proceedings of the 2001 International Conference on Compilers, Architecture, and Synthesis for Embedded Systems*, CASES '01, pages 238–246, New York, NY, USA, 2001. ACM.
- [13] Aki Saarinen, Matti Siekkinen, Yu Xiao, Jukka K. Nurminen, Matti Kempainen, and Pan Hui. Can offloading save energy for popular apps? In *Proceedings of the Seventh ACM International Workshop on Mobility in the Evolving Internet Architecture*, MobiArch '12, pages 3–10, New York, NY, USA, 2012. ACM.
- [14] M. Satyanarayanan. A brief history of cloud offload. *GetMobile*, 18(4):19–23, Oct 2014.
- [15] Mahadev Satyanarayanan, P. Bahl, R. Caceres, and N. Davies. The case for vm-based cloudlets in mobile computing. *Pervasive Computing, IEEE*, 8(4):14–23, Oct 2009.